

HicranArnold_project_data_process

December 12, 2022

Auther : Hicran Arnold

MET CS 6999 Term Project

Model Creations- Iterations

0.1 Libraries Used

```
[ ]: # libraries
# for data processing
import pandas as pd # for reading and analyzing the data
#for visualizations
import seaborn as sns
import matplotlib.pyplot as plt
from pandas_profiling import ProfileReport
#to split data into training and testing
from sklearn.model_selection import train_test_split
##feature selection
from sklearn.feature_selection import mutual_info_classif # attribute selection
    ↳method I
from sklearn.feature_selection import chi2 # attribute selection method II
from sklearn.ensemble import RandomForestClassifier # attribute selection
    ↳method III
from sklearn.feature_selection import SelectFromModel # for attribute selection
    ↳method III
from mlxtend.feature_selection import SequentialFeatureSelector as SFS # for
    ↳attribute selection method IV
##classifications
from sklearn.linear_model import LogisticRegression #clf1
from sklearn.tree import DecisionTreeClassifier #clf2
from sklearn.neighbors import KNeighborsClassifier#clf3
from sklearn import svm #clf4
from sklearn.neural_network import MLPClassifier#clf5
#accuracy measurement
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import log_loss, roc_auc_score, recall_score,
    ↳precision_score, average_precision_score, f1_score, classification_report
from sklearn import metrics
```

```
#data transformation
from sklearn import preprocessing
```

0.2 Reading and Splitting the Data

```
[ ]: # saving as df
df= pd.read_csv("HicranArnold_project_intial.csv", index_col = False)
df.head()
```

```
[ ]:      x.aidtst3  employ1  income2  weight2  height3  children  veteran3  blind  \
0          2.0        1.0        8.0    138.0    504.0         2.0         2.0    2.0
1          1.0        1.0        4.0    240.0    507.0        88.0         2.0    1.0
2          1.0        1.0        6.0    120.0    504.0        88.0         2.0    2.0
3          2.0        7.0       77.0    133.0    502.0        88.0         2.0    2.0
4          9.0        7.0        8.0    170.0    505.0        88.0         2.0    2.0
```

```
      renthom1  sex1  ...  x.denvst3  x.prace1  x.mrace1  x.exteth3  x.asthms1  \
0          1.0    2  ...           1          1          1          1          3
1          2.0    1  ...           2          1          1          2          3
2          1.0    2  ...           1          6          6          2          3
3          1.0    2  ...           1          1          1          2          3
4          1.0    2  ...           1          1          1          1          3
```

```
      x.michd  x.ltasth1  x.casthm1  x.state  havarth3
0          2.0          1          1       27         2
1          2.0          1          1       72         2
2          2.0          1          1       31         2
3          2.0          1          1       45         2
4          2.0          1          1       24         2
```

[5 rows x 108 columns]

```
[ ]: X = df.loc[ :, ~df.columns.isin(['havarth3']) ]
y = df.loc[:, 'havarth3':'havarth3']
```

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
↪, random_state = 42)
print("original data size ", df.shape)
print("training data size", X_train.shape)
print("testing data size ", X_test.shape)
# print(X_train.head())
```

```
original data size (11933, 108)
training data size (7995, 107)
testing data size (3938, 107)
```

0.3 Saving Data Training and Test CSV Files

```
[ ]: X_train_csv = X_train.copy()
X_train_csv["havarth3"] = y_train.copy()
X_train_csv.to_csv("HicranArnold_project_training.csv")
```

```
[ ]: X_test_csv = X_test.copy()
X_test_csv["havarth3"] = y_test.copy()
X_test_csv.to_csv("HicranArnold_project_test.csv")
```

```
[ ]: print("training data size", X_train.shape)
print("testing y data size ", y_train.shape)
print("testing data size ", X_test.shape)
print("testing y data size ", y_test.shape)
```

```
training data size (7995, 107)
testing y data size (7995, 1)
testing data size (3938, 107)
testing y data size (3938, 1)
```

0.4 Identifying Numeric and Categorical Features

```
[ ]: import numpy as np
categ_data = []
numeric_data = []
X_train_col = X_train.columns.tolist()
for column in X_train_col:
    unique_values = len(X_train[column].value_counts())
    if unique_values < 10:
        categ_data.append(column)
    else:
        numeric_data.append(column)
print("categorical data")
print(categ_data)
print("numeric data")
print(numeric_data)
```

```
categorical data
['x.aidtst3', 'employ1', 'veteran3', 'blind', 'renthom1', 'sex1', 'marital',
'educa', 'deaf', 'decide', 'x.drnkdrv', 'flushot6', 'seatbelt', 'hivtst6',
'hivrisk5', 'pneuvac4', 'diffwalk', 'usenow3', 'diffdres', 'diffalon',
'smoke100', 'rmvteth4', 'lastden4', 'diabete3', 'hlthpln1', 'genhlth',
'dispcode', 'chckdny1', 'iyear', 'persdoc2', 'medcost', 'checkup1', 'exerany2',
'chcocncr', 'chccopd1', 'addepev2', 'chcscncr', 'asthma3', 'cvdstrk3',
'cvdinfr4', 'cvdcrhd4', 'qstver', 'qstlang', 'x.metstat', 'x.bmi5cat',
'x.age.g', 'x.raceg21', 'x.race.g1', 'x.age65yr', 'x.rfbmi5', 'x.chldcnt',
'x.educag', 'x.incomg', 'x.rfdrhv6', 'x.rfseat2', 'x.rfseat3', 'x.rfbing5',
'x.smoker3', 'x.rfsmok3', 'drnkany5', 'x.racegr3', 'x.race', 'x.urbstat',
```

```
'x.chispnc', 'x.rfhlth', 'x.dualuse', 'x.imprace', 'x.hispanc', 'x.phys14d',
'x.ment14d', 'x.hcvu651', 'x.totinda', 'x.denvst3', 'x.prace1', 'x.mrace1',
'x.exteth3', 'x.asthms1', 'x.michd', 'x.ltasth1', 'x.casthm1']
numeric data
['income2', 'weight2', 'height3', 'children', 'alcdays5', 'x.psu', 'phys14d',
'menth14d', 'fmonth', 'imonth', 'iday', 'sleptim1', 'htin4', 'wtkg3', 'x.bmi5',
'htm4', 'x.age80', 'x.ageg5yr', 'x.drnkwek', 'drocdy3.', 'x.llcpwt2',
'x.llcpwt', 'x.wt2rake', 'x.ststr', 'x.strwt', 'x.rawrake', 'x.state']
```

```
[ ]: numeric_var = ["weight2","height3","children","alcdays5","x.
↳psu","phys14d","ment14d","sleptim1","htin4","wtkg3","x.bmi5","htm4","x.
↳drnkwek","drocdy3.","x.llcpwt2","x.llcpwt","x.wt2rake","x.ststr"
,"x.strwt","x.rawrake"]

categorical = X_train.loc[:, ~X_train.columns.isin(numeric_var)].columns.
↳to_list()
```

```
[ ]: list(set(numeric_data) - set(numeric_var))
```

```
[ ]: ['x.age80', 'imonth', 'x.ageg5yr', 'income2', 'iday', 'x.state', 'fmonth']
```

We have have only above differences. We know that dates are not count as numeric and ranges, state number not a numeric variable. We are now sure that our data is correct.

0.5 Iteration One

0.5.1 Attribute Selection Method : Information Gain/Mutual Information

First attribute selection method is Information Gain/Mutual Information. This is a filtering method. It means we will use univariate statistics to pick our attributes and filter them based on their information gain. We will only select the ones with high information gain.

```
[ ]: X_train.columns
```

```
[ ]: Index(['x.aidtst3', 'employ1', 'income2', 'weight2', 'height3', 'children',
'veteran3', 'blind', 'renthom1', 'sex1',
...
'x.totinda', 'x.denvst3', 'x.prace1', 'x.mrace1', 'x.exteth3',
'x.asthms1', 'x.michd', 'x.ltasth1', 'x.casthm1', 'x.state'],
dtype='object', length=107)
```

```
[ ]: # takes df, return to df
def mutualInfo(X_train, y_train, X_test):
    #initating the mututal info object
    mutual_info = mutual_info_classif(X_train, y_train.values.ravel(),
↳random_state=42)
    #coverting to series
    mutual_info = pd.Series(mutual_info)
```

```

#renaming index with column names
mutual_info.index = X_train.columns
#sorting
mutual_info.sort_values(ascending=False)
#mutual_info.between(0.1,1 ) because we do not want the one that has zero
↪rate
#filtering the data with only selected attributes
mutual_info[mutual_info.between(0.01,1 )].sort_values(ascending=False).plot.
↪bar(figsize=(8, 8))
X_train_reduced= X_train[mutual_info[mutual_info.between(0.01,1 )].
↪sort_values(ascending=False)[0:20].index.to_list()]
X_test_reduced = X_test[mutual_info[mutual_info.between(0.01,1 )].
↪sort_values(ascending=False)[0:20].index.to_list()]
return X_train_reduced, X_test_reduced

```

```


[ ]: # reduced data set 1
X_train_reduced1 , X_test_reduced1= mutualInfo(X_train, y_train, X_test)
print("new size X_train",X_train_reduced1.shape )
print("new size X_train",X_test_reduced1.shape )
print("selected column names",X_test_reduced1.columns)

```

```

new size X_train (7995, 20)
new size X_train (3938, 20)
selected column names Index(['x.age80', 'x.ageg5yr', 'x.age.g', 'employ1',
'diffwalk', 'x.hcvu651',
'genhlth', 'x.age65yr', 'marital', 'x.phys14d', 'physhlth', 'rmvteth4',
'pneuvac4', 'x.rfhlth', 'x.ststr', 'x.exteth3', 'persdoc2', 'qstver',
'x.strwt', 'chccopd1'],
dtype='object')

```



0.5.2 Model Creation : Initializing 5 models

Model 1: Logistic Regression Classifier 1- Logistic Regression

```
[ ]: def lRMaker(X_train, y_train, X_test):
    # instantiate the model (using the default parameters)
    m = LogisticRegression (solver='lbfgs', max_iter=1000, random_state=16) #
    ↪ initialize lr
    m.fit(X_train, y_train.values.ravel()) # fit the training data to train the
    ↪ data
    y_pred = m.predict(X_test) # use test data to predict
```

```

print(y_pred)
return m, y_pred

```

```

[ ]: def printAccuracy(y_test, y_pred):
    print(y_pred)
    cm = confusion_matrix(y_test, y_pred)
    TP, FP, FN, TN = cm.ravel()
    group_names = ["TP", "FP", "FN", "TN"]
    group_counts = ["{0:0.0f}".format(value) for value in
                    cm.flatten()]
    group_percentages = ["{0:.2%}".format(value) for value in
                        cm.flatten()/np.sum(cm)]
    labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in
              zip(group_names, group_counts, group_percentages)]
    labels = np.asarray(labels).reshape(2,2)
    sns.heatmap(cm, annot=labels, fmt='', cmap='mako')
    P = TP+FP
    N= FN+ TN
    TPR_log = TP/P
    TNR_log = TN /N

    my_accuracy_score_log= accuracy_score(y_test, y_pred)
    my_matrix_values_1 = {"TN":TN, "FP":FP, "FN":FN, "TP":TP, "Accuracy":
↪my_accuracy_score_log,
                        "Sensitivity: TPR":TPR_log, "Specificity: TNR":TNR_log}
    result_of_m = pd.DataFrame(my_matrix_values_1 , index=[0])
    print(result_of_m)
    #return my_accuracy_score_log
    print('Log loss = {:.5f}'.format(log_loss(y_test, y_pred)))
    print('AUC = {:.5f}'.format(roc_auc_score(y_test, y_pred)))
    print('Average Precision = {:.5f}'.
↪format(average_precision_score(y_test, y_pred)))
    print('\nUsing 0.5 as threshold:')
    print('Accuracy = {:.5f}'.format(accuracy_score(y_test, y_pred)))
    print('Precision = {:.5f}'.format(precision_score(y_test, y_pred)))
    print('Recall = {:.5f}'.format(recall_score(y_test, y_pred)))
    print('F1 score = {:.5f}'.format(f1_score(y_test, y_pred)))
    print('\nClassification Report')
    print(classification_report(y_test, y_pred))
    print("Note that in binary classification, recall of the positive class_
↪is also known as \"sensitivity\"; recall of the negative class is_
↪\"specificity\".")
    return result_of_m

```

```

[ ]: m1_lr, m1_y_pred_lr= lRMaker(X_train=X_train_reduced1, y_train=y_train,
↪X_test=X_test_reduced1)

```

```
[1 2 1 ... 2 2 1]
```

```
[ ]: m1_lr_result = printAccuracy(y_test, m1_y_pred_lr)
```

```
[1 2 1 ... 2 2 1]
```

```
      TN   FP   FN   TP  Accuracy  Sensitivity: TPR  Specificity: TNR
0  2212  785  418  523  0.694515          0.399847          0.841065
```

```
Log loss = 11.97184
```

```
AUC = 0.62046
```

```
Average Precision = 0.29001
```

```
Using 0.5 as threshold:
```

```
Accuracy = 0.69451
```

```
Precision = 0.55579
```


```
Recall = 0.39985
```

```
F1 score = 0.46510
```

```
Classification Report
```

	precision	recall	f1-score	support
1	0.56	0.40	0.47	1308
2	0.74	0.84	0.79	2630
accuracy			0.69	3938
macro avg	0.65	0.62	0.63	3938
weighted avg	0.68	0.69	0.68	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 2:Decision Tree Classifier *Classifier 2-Decision Tree Classifier*

```
[ ]: def dTMaker(X_train, y_train, X_test):
    #max_depth to control the size of the tree to prevent overfitting
    m = DecisionTreeClassifier(max_depth=3, random_state=42)
    m = m.fit(X_train, y_train.values.ravel())
    y_pred = m.predict(X_test)
    return m , y_pred

[ ]: m2_dt, m2_y_pred_dt = dTMaker (X_train=X_train_reduced1, y_train=y_train,
    ↪X_test=X_test_reduced1)

[ ]: # printing dt result
m2_dt_result = printAccuracy(y_test, m2_y_pred_dt)
```

```
[2 2 1 ... 2 2 1]
      TN   FP   FN   TP  Accuracy  Sensitivity: TPR  Specificity: TNR
0  2276  652  354  656   0.74454      0.501529      0.865399
Log loss = 11.97184
AUC = 0.68346
Average Precision = 0.27758
```

Using 0.5 as threshold:

Accuracy = 0.74454

Precision = 0.64950


Recall = 0.50153

F1 score = 0.56601

Classification Report

	precision	recall	f1-score	support
1	0.65	0.50	0.57	1308
2	0.78	0.87	0.82	2630
accuracy			0.74	3938
macro avg	0.71	0.68	0.69	3938
weighted avg	0.73	0.74	0.73	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 3 : K-Nearest Neighbors Classifier (KNN)

```
[ ]: def kNNMaker(X_train,y_train,X_test):
    scaler =preprocessing.StandardScaler().fit(X_train)
    X_train_scaled = scaler.transform(X_train)
    X_test_scaled= scaler.transform(X_test)
    m = KNeighborsClassifier(n_neighbors=9)
    m.fit(X_train_scaled, y_train.values.ravel())
    y_pred_knn= m.predict(X_test_scaled)
    return m , y_pred_knn

[ ]: m3_knn, m3_y_pred_knn = kNNMaker(X_train=X_train_reduced1, y_train=y_train,
↪X_test=X_test_reduced1)

[ ]: m3_knn_result = printAccuracy(y_test, m3_y_pred_knn)
```

```
[1 2 1 ... 2 1 1]
      TN   FP   FN   TP Accuracy Sensitivity: TPR Specificity: TNR
0  2203  627  427  681  0.732351          0.520642          0.837643
Log loss = 11.97184
AUC = 0.67914
Average Precision = 0.27913
```


Using 0.5 as threshold:

```
Accuracy = 0.73235
Precision = 0.61462
Recall = 0.52064
F1 score = 0.56374
```

Classification Report

	precision	recall	f1-score	support
1	0.61	0.52	0.56	1308
2	0.78	0.84	0.81	2630
accuracy			0.73	3938
macro avg	0.70	0.68	0.69	3938
weighted avg	0.72	0.73	0.73	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 4 :Support vector machine (SVM) *Classifier 4 : Support vector machine (SVM)*

```
[ ]: def sVMMaker(X_train,y_train,X_test ):
    scaler =preprocessing.StandardScaler().fit(X_train)
    X_train_scaled = scaler.transform(X_train)
    X_test_scaled= scaler.transform(X_test)
    m= svm.SVC(random_state=0)
    m = m.fit(X_train_scaled, y_train.values.ravel())
    y_pred_svm= m.predict(X_test_scaled)
    return m, y_pred_svm

[ ]: m4_svm, m4_y_pred_svm = sVMMaker(X_train=X_train_reduced1, y_train=y_train,
    ↪X_test=X_test_reduced1)

[ ]: m4_svm_result = printAccuracy(y_test, m4_y_pred_svm)
```

```
[2 2 1 ... 2 2 1]
      TN   FP   FN   TP  Accuracy  Sensitivity: TPR  Specificity: TNR
0  2351  690  279  618  0.753936      0.472477      0.893916
Log loss = 11.97184
AUC = 0.68320
```

Average Precision = 0.27663

Using 0.5 as threshold:

Accuracy = 0.75394

Precision = 0.68896


Recall = 0.47248

F1 score = 0.56054

Classification Report

	precision	recall	f1-score	support
1	0.69	0.47	0.56	1308
2	0.77	0.89	0.83	2630
accuracy			0.75	3938
macro avg	0.73	0.68	0.69	3938
weighted avg	0.75	0.75	0.74	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 5 neural_network.MLPClassifie *Classifier 5: neural_network.MLPClassifier*

```
[ ]: def mLPMaker(X_train,X_test, y_train):
    scaler =preprocessing.StandardScaler().fit(X_train)
    X_train_scaled = scaler.transform(X_train)
    X_test_scaled= scaler.transform(X_test)
    m = MLPClassifier(random_state=0, max_iter=1).fit(X_train_scaled, y_train.
↪values.ravel())
    y_pred_mlp= m.predict(X_test_scaled)
    return m, y_pred_mlp
```

```
[ ]: m5_mlp , m5_y_pred_mlp = mLPMaker(X_train=X_train_reduced1, y_train=y_train,↪
↪X_test=X_test_reduced1)
```

/home/harnold/github/2018_BRFSS_survey_data_anlaysiais/env/lib/python3.8/site-packages/sklearn/neural_network/_multilayer_perceptron.py:679:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1) reached and the optimization hasn't converged yet.
warnings.warn(

```
[ ]: m5_mlp_result = printAccuracy(y_test, m5_y_pred_mlp)
```

```
[2 2 1 ... 2 2 1]
      TN   FP   FN   TP Accuracy Sensitivity: TPR Specificity: TNR
0  2302  681  328  627  0.743779          0.479358          0.875285
Log loss = 11.97184
AUC = 0.67732
Average Precision = 0.27808
```


Using 0.5 as threshold:

```
Accuracy = 0.74378
Precision = 0.65654
Recall = 0.47936
F1 score = 0.55413
```

Classification Report

	precision	recall	f1-score	support
1	0.66	0.48	0.55	1308
2	0.77	0.88	0.82	2630
accuracy			0.74	3938
macro avg	0.71	0.68	0.69	3938
weighted avg	0.73	0.74	0.73	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



0.6 Iteration 2

0.6.1 Attribute Selection Method : Chi-square Test

```
[ ]: def chiFeatures(X_train,numeric_var, y_train):
    X_categorical= X_train.loc[ :, ~X_train.columns.isin(numeric_var)]
    f_p_values=chi2(X_categorical,y_train)
    # higher the p value lower the importance
    # A p-value less than 0.05 is typically considered to be statistically
    ↪significant
    p_values=pd.Series(f_p_values[1])
    p_values.index=X_categorical.columns
    p_values.sort_values(ascending=False, inplace=True)
    column_name= p_values.index.to_list()
    values= p_values.values.tolist()
    data= pd.DataFrame(data=values,columns=['p_values'] ) # 0000001
    data["featureNames"] = column_name
    # this will eliminate our features to 20
    cat= data[data["p_values"]>0.000000000000000000000000000001].featureNames.
    ↪to_list()
    len(cat)
    return cat
```

```
def dropCorr(categorical_var, X_train,X_test ):
    # categorical elimination
    X_train_reduced= X_train.loc[ :, ~X_train.columns.isin(categorical_var) ]
    X_test_reduced = X_test.loc[ :, ~X_test.columns.isin(categorical_var) ]
    return X_train_reduced, X_test_reduced
```

```
[ ]: categorical_var = chiFeatures(X_train,numeric_var, y_train)
X_train_reduced2, X_test_reduced2= dropCorr(categorical_var, X_train,X_test )
print("new size X_train",X_train_reduced2.shape )
print("new size X_test",X_test_reduced2.shape )
print("selected column names",X_train_reduced2.columns)
```

```
new size X_train (7995, 37)
new size X_test (3938, 37)
selected column names Index(['employ1', 'income2', 'weight2', 'height3',
'children', 'x.drnkdrv',
'pneuvac4', 'alcdays5', 'rmvteth4', 'x.psu', 'physhlth', 'menthlth',
'genhlth', 'sleptim1', 'qstver', 'htin4', 'wtkg3', 'x.bmi5', 'htm4',
'x.age.g', 'x.age80', 'x.age5yr', 'x.chldcnt', 'x.drnkwek', 'drocdy3.',
'x.race', 'x.llcpwt2', 'x.llcpwt', 'x.dualuse', 'x.wt2rake', 'x.ststr',
'x.strwt', 'x.rawrake', 'x.phys14d', 'x.hcvu651', 'x.prace1',
'x.mrace1'],
dtype='object')
```

0.6.2 Model Creation : Initializing 5 models

Model 6 : Logistic Regression

```
[ ]: m6_lr, m6_y_pred_lr= lRMaker(X_train=X_train_reduced2, y_train=y_train,
↪X_test=X_test_reduced2)
```

```
[2 2 2 ... 2 2 2]
```

```
[ ]: m6_lr_result = printAccuracy(y_test, m6_y_pred_lr)
```

```
[2 2 2 ... 2 2 2]
```

```
      TN      FP  FN  TP  Accuracy  Sensitivity: TPR  Specificity: TNR
0  2630  1308   0   0  0.667852                0.0                1.0
Log loss = 11.97184
AUC = 0.50000
Average Precision = 0.33215
```

Using 0.5 as threshold:

```
Accuracy = 0.66785
Precision = 0.00000
Recall = 0.00000
F1 score = 0.00000
```


Classification Report

	precision	recall	f1-score	support
1	0.00	0.00	0.00	1308
2	0.67	1.00	0.80	2630
accuracy			0.67	3938
macro avg	0.33	0.50	0.40	3938
weighted avg	0.45	0.67	0.53	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.

/home/harnold/github/2018_BRFSS_survey_data_anlaysiais/env/lib/python3.8/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

/home/harnold/github/2018_BRFSS_survey_data_anlaysiais/env/lib/python3.8/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.


_warn_prf(average, modifier, msg_start, len(result))

/home/harnold/github/2018_BRFSS_survey_data_anlaysiais/env/lib/python3.8/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

/home/harnold/github/2018_BRFSS_survey_data_anlaysiais/env/lib/python3.8/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))



Model 7 : Decision Tree

```
[ ]: m7_dt, m7_y_pred_dt = dTMaker (X_train=X_train_reduced2, y_train=y_train,
    ↪X_test=X_test_reduced2)
```

```
[ ]: m7_dt_result = printAccuracy(y_test, m7_y_pred_dt)
```

```
[2 2 1 ... 2 2 1]
      TN  FP  FN  TP Accuracy Sensitivity: TPR Specificity: TNR
0  2423  871  207  437  0.726257      0.334098      0.921293
Log loss = 11.97184
AUC = 0.62770
Average Precision = 0.28705
```

Using 0.5 as threshold:


```
Accuracy = 0.72626
Precision = 0.67857
Recall = 0.33410
F1 score = 0.44775
```

Classification Report

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

1	0.68	0.33	0.45	1308
2	0.74	0.92	0.82	2630
accuracy			0.73	3938
macro avg	0.71	0.63	0.63	3938
weighted avg	0.72	0.73	0.70	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 8: KNN

```
[ ]: m8_knn, m8_y_pred_knn = kNNMaker(X_train=X_train_reduced2, y_train=y_train,
↪X_test=X_test_reduced2)
```

```
[ ]: m8_knn_result = printAccuracy(y_test, m8_y_pred_knn)
```

```
[2 2 1 ... 2 2 1]
      TN  FP  FN  TP  Accuracy  Sensitivity: TPR  Specificity: TNR
0  2201  676  429  632  0.719401          0.48318          0.836882
Log loss = 11.97184
AUC = 0.66003
Average Precision = 0.28192
```

Using 0.5 as threshold:

Accuracy = 0.71940

Precision = 0.59566


Recall = 0.48318

F1 score = 0.53356

Classification Report

	precision	recall	f1-score	support
1	0.60	0.48	0.53	1308
2	0.77	0.84	0.80	2630
accuracy			0.72	3938
macro avg	0.68	0.66	0.67	3938
weighted avg	0.71	0.72	0.71	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 9: SVP

```
[ ]: m9_svm, m9_y_pred_svm = sVMMaker(X_train=X_train_reduced2, y_train=y_train,
↪X_test=X_test_reduced2)
```

```
[ ]: m9_svm_result = printAccuracy(y_test, m9_y_pred_svm)
```

```
[2 2 1 ... 2 2 1]
```

	TN	FP	FN	TP	Accuracy	Sensitivity: TPR	Specificity: TNR
0	2325	693	305	615	0.746572	0.470183	0.88403

```
Log loss = 11.97184
```

```
AUC = 0.67711
```

```
Average Precision = 0.27783
```

```
Using 0.5 as threshold:
```

```
Accuracy = 0.74657
```

```
Precision = 0.66848
```


```
Recall = 0.47018
```

```
F1 score = 0.55206
```

```
Classification Report
```

	precision	recall	f1-score	support
1	0.67	0.47	0.55	1308
2	0.77	0.88	0.82	2630
accuracy			0.75	3938
macro avg	0.72	0.68	0.69	3938
weighted avg	0.74	0.75	0.73	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 10: MLP

```
[ ]: m10_mlp , m10_y_pred_mlp = mLPMaker(X_train=X_train_reduced2, y_train=y_train,
    ↪X_test=X_test_reduced2)
```

```
/home/harnold/github/2018_BRFSS_survey_data_anlalysis/env/lib/python3.8/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:679:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1) reached and the
optimization hasn't converged yet.
    warnings.warn(
```

```
[ ]: m10_dt_result = printAccuracy(y_test, m10_y_pred_mlp)
```

```
[2 2 1 ... 2 2 1]
```

	TN	FP	FN	TP	Accuracy	Sensitivity: TPR	Specificity: TNR
0	2274	654	356	654	0.743525	0.5	0.864639

```
Log loss = 11.97184
```

```
AUC = 0.68232
```

```
Average Precision = 0.27775
```

```
Using 0.5 as threshold:
```

```
Accuracy = 0.74352
```


```
Precision = 0.64752
```

Recall = 0.50000
F1 score = 0.56428

Classification Report

	precision	recall	f1-score	support
1	0.65	0.50	0.56	1308
2	0.78	0.86	0.82	2630
accuracy			0.74	3938
macro avg	0.71	0.68	0.69	3938
weighted avg	0.73	0.74	0.73	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



0.7 Iteration 3

0.7.1 Attribute Selection Method : Random Forest Classifier

```
[ ]: def randomForestSelect(X_train,X_test,y_train, max_features):  
    sel = SelectFromModel(RandomForestClassifier(n_estimators = 100,  
↳random_state=0,max_features=20))  
    sel.fit(X_train, y_train.values.ravel())  
    selected_feat= X_train.columns[(sel.get_support())]  
    len(selected_feat)  
    #print(selected_feat)  
    X_train_selected = X_train[selected_feat]  
    X_test_selected = X_test[selected_feat]  
    print(X_test_selected.head())  
    return X_train_selected, X_test_selected
```

```
[ ]: # reduced data set 3  
X_train_reduced3 , X_test_reduced3=↳  
↳randomForestSelect(X_train=X_train,X_test=X_test,y_train=↳  
↳y_train,max_features=20)  
print("new size X_train",X_train_reduced3.shape )  
print("new size X_test",X_test_reduced3.shape )  
print("selected column names",X_train_reduced3.columns)
```

	employ1	income2	weight2	height3	pneuvac4	alcdays5	diffwalk	\
10172	5.0	7.0	130.0	502.0	1.0	101.0	2.0	
10322	1.0	7.0	170.0	508.0	2.0	888.0	2.0	
9195	7.0	99.0	145.0	501.0	1.0	215.0	2.0	
357	2.0	7.0	200.0	511.0	2.0	201.0	2.0	
4352	1.0	8.0	180.0	507.0	2.0	888.0	2.0	

	x.psu	physlth	menthlth	...	x.incomg	x.drnkwek	drocdy3.	\
10172	2018001646	88	88	...	5	99900	14	
10322	2018002740	88	88	...	5	0	0	
9195	2018000479	30	88	...	9	350	50	
357	2018004168	88	88	...	5	47	3	
4352	2018004398	88	88	...	9	99900	900	

	x.llcpwt2	x.llcpwt	x.wt2rake	x.ststr	x.strwt	x.phys14d	\
10172	462.325198	327.156320	36.202567	51031	18.101283	1	
10322	162.085147	274.500192	39.744340	201041	13.248113	1	
9195	268.171190	166.046135	21.603358	271031	21.603358	3	
357	1093.379630	2355.282136	56.130293	132019	56.130293	1	
4352	158.219742	83.045951	28.192596	442019	28.192596	1	

	x.state
10172	5
10322	20


```

9195      27
357       13
4352      44

```

```

[5 rows x 34 columns]
new size X_train (7995, 34)
new size X_train (3938, 34)
selected column names Index(['employ1', 'income2', 'weight2', 'height3',
'pneuvac4', 'alcdays5',
'diffwalk', 'x.psu', 'physhlth', 'menthlth', 'genhlth', 'fmonth',
'imonth', 'iday', 'sleptim1', 'qstver', 'htin4', 'wtkg3', 'x.bmi5',
'htm4', 'x.age.g', 'x.age80', 'x.age5yr', 'x.age65yr', 'x.incomg',
'x.drnkwek', 'drocdy3.', 'x.llcpwt2', 'x.llcpwt', 'x.wt2rake',
'x.ststr', 'x.strwt', 'x.phys14d', 'x.state'],
dtype='object')

```

0.7.2 Model Creation : Initializing 5 models

Model 11 : Logistic Regression

```
[ ]: m11_lr, m11_y_pred_lr= lRMaker(X_train=X_train_reduced3, y_train=y_train,
↪X_test=X_test_reduced3)
```

```
[2 2 2 ... 2 2 2]
```

```
[ ]: m11_lr_result = printAccuracy(y_test, m11_y_pred_lr)
```

```
[2 2 2 ... 2 2 2]
```

```

      TN      FP  FN  TP  Accuracy  Sensitivity: TPR  Specificity: TNR
0  2630  1308   0   0  0.667852              0.0              1.0
Log loss = 11.97184
AUC = 0.50000
Average Precision = 0.33215

```

Using 0.5 as threshold:

```

Accuracy = 0.66785
Precision = 0.00000
Recall = 0.00000
F1 score = 0.00000

```

Classification Report

	precision	recall	f1-score	support
1	0.00	0.00	0.00	1308
2	0.67	1.00	0.80	2630
accuracy			0.67	3938
macro avg	0.33	0.50	0.40	3938
weighted avg	0.45	0.67	0.53	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.

```
/home/harnold/github/2018_BRFSS_survey_data_anlaysiais/env/lib/python3.8/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/home/harnold/github/2018_BRFSS_survey_data_anlaysiais/env/lib/python3.8/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```


```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/home/harnold/github/2018_BRFSS_survey_data_anlaysiais/env/lib/python3.8/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/home/harnold/github/2018_BRFSS_survey_data_anlaysiais/env/lib/python3.8/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```



Model 12: Decision Tree

```
[ ]: m12_dt, m12_y_pred_dt = dTMaker(X_train=X_train_reduced3, y_train=y_train,  
    ↪X_test=X_test_reduced3)
```

```
[ ]: m12_dt_result = printAccuracy(y_test, m12_y_pred_dt)
```

```
[2 2 1 ... 2 2 1]
```

	TN	FP	FN	TP	Accuracy	Sensitivity: TPR	Specificity: TNR
0	2275	656	355	652	0.743271	0.498471	0.865019

Log loss = 11.97184

AUC = 0.68174

Average Precision = 0.27782

Using 0.5 as threshold:

Accuracy = 0.74327

Precision = 0.64747


Recall = 0.49847

F1 score = 0.56328

Classification Report

	precision	recall	f1-score	support
1	0.65	0.50	0.56	1308
2	0.78	0.87	0.82	2630
accuracy			0.74	3938
macro avg	0.71	0.68	0.69	3938
weighted avg	0.73	0.74	0.73	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 13: KNN

```
[ ]: m13_knn, m13_y_pred_knn = kNNMaker(X_train=X_train_reduced3, y_train=y_train,
    ↪X_test=X_test_reduced3)
```

```
[ ]: m13_knn_result = printAccuracy(y_test, m13_y_pred_knn)
```

```
[2 2 1 ... 2 2 1]
```

	TN	FP	FN	TP	Accuracy	Sensitivity: TPR	Specificity: TNR
0	2237	667	393	641	0.730828	0.490061	0.85057

Log loss = 11.97184

AUC = 0.67032

Average Precision = 0.27990

Using 0.5 as threshold:

Accuracy = 0.73083

Precision = 0.61992

Recall = 0.49006


F1 score = 0.54740

Classification Report

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

1	0.62	0.49	0.55	1308
2	0.77	0.85	0.81	2630
accuracy			0.73	3938
macro avg	0.70	0.67	0.68	3938
weighted avg	0.72	0.73	0.72	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 14: SVM

```
[ ]: m14_svm, m14_y_pred_svm = sVMMaker(X_train=X_train_reduced3, y_train=y_train,
    ↪X_test=X_test_reduced3)
```

```
[ ]: m14_svm_result = printAccuracy(y_test, m14_y_pred_svm)
```

```
[2 2 1 ... 2 2 1]
      TN  FP  FN  TP  Accuracy  Sensitivity: TPR  Specificity: TNR
0  2347  694  283  614  0.751905      0.469419      0.892395
Log loss = 11.97184
AUC = 0.68091
Average Precision = 0.27700
```

Using 0.5 as threshold:

Accuracy = 0.75190

Precision = 0.68450


Recall = 0.46942

F1 score = 0.55692

Classification Report

	precision	recall	f1-score	support
1	0.68	0.47	0.56	1308
2	0.77	0.89	0.83	2630
accuracy			0.75	3938
macro avg	0.73	0.68	0.69	3938
weighted avg	0.74	0.75	0.74	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 15: MLP

```
[ ]: m15_mlp , m15_y_pred_mlp = mLPMaker(X_train=X_train_reduced3, y_train=y_train,
↪X_test=X_test_reduced3)
```

```
/home/harnold/github/2018_BRFSS_survey_data_anlaysia/env/lib/python3.8/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:679:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1) reached and the
optimization hasn't converged yet.
warnings.warn(
```

```
[ ]: m15_mlp_result = printAccuracy(y_test, m15_y_pred_mlp)
```

```
[2 2 1 ... 2 2 1]
      TN   FP   FN   TP Accuracy Sensitivity: TPR Specificity: TNR
0  2273  664  357  644  0.740731          0.492355          0.864259
Log loss = 11.97184
AUC = 0.67831
Average Precision = 0.27830
```


```
Using 0.5 as threshold:
Accuracy = 0.74073
Precision = 0.64336
Recall = 0.49235
F1 score = 0.55782
```

```
Classification Report
              precision    recall  f1-score   support

         1              0.64       0.49       0.56       1308
         2              0.77       0.86       0.82       2630

 accuracy              0.74       0.74       0.74       3938
 macro avg              0.71       0.68       0.69       3938
weighted avg              0.73       0.74       0.73       3938
```

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



0.8 Iteration 4

0.8.1 Attribute selection method: Forward Feature Selection, this method is a cost expensive method and it takes a while to compute therefore we will reduce the feature amount to 7

This model did not selected good pairs of attributes

```
[ ]: def sfsFeatures(X_train,X_test, y_train):
    model = LogisticRegression(solver='lbfgs', max_iter=1000, random_state=0)
    sfs1 = SFS(model, k_features=7, forward=True, floating=False,
    verbose=2,scoring='accuracy',n_jobs=-1,cv=5)
    sfs1 = sfs1.fit(X_train, y_train.values.ravel())
    features = list(sfs1.k_feature_names_)
    X_train_reduced =X_train[features]
    X_test_reduced= X_test[features]
    #pd.DataFrame.from_dict(sfs1.get_metric_dict()).T
    return X_train_reduced, X_test_reduced
```

```
[ ]: # reduced data set 4
X_train_reduced4 , X_test_reduced4=
    sfsFeatures(X_train=X_train,X_test=X_test,y_train= y_train)
```



```
print("new size X_train",X_train_reduced4.shape )
print("new size X_test",X_test_reduced4.shape )
print("selected column names",X_train_reduced4.columns)
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    1.2s
[Parallel(n_jobs=-1)]: Done 107 out of 107 | elapsed:    1.7s finished
```

```
[2022-12-12 21:06:17] Features: 1/7 -- score:
0.7292057535959975[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8
concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:    0.5s
[Parallel(n_jobs=-1)]: Done 106 out of 106 | elapsed:    1.0s finished
```

```
[2022-12-12 21:06:18] Features: 2/7 -- score:
0.7295809881175735[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8
concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:    0.5s
[Parallel(n_jobs=-1)]: Done 105 out of 105 | elapsed:    1.2s finished
```

```
[2022-12-12 21:06:19] Features: 3/7 -- score:
0.7307066916823015[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8
concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:    0.9s
[Parallel(n_jobs=-1)]: Done 89 out of 104 | elapsed:    2.1s remaining:    0.4s
[Parallel(n_jobs=-1)]: Done 104 out of 104 | elapsed:    2.3s finished
```

```
[2022-12-12 21:06:22] Features: 4/7 -- score:
0.7310819262038775[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8
concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:    1.1s
[Parallel(n_jobs=-1)]: Done 103 out of 103 | elapsed:    2.9s finished
```

```
[2022-12-12 21:06:25] Features: 5/7 -- score:
0.7313320825515949[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8
concurrent workers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    1.2s
[Parallel(n_jobs=-1)]: Done 102 out of 102 | elapsed:    4.6s finished
```

```
[2022-12-12 21:06:29] Features: 6/7 -- score:
0.7317073170731707[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8
concurrent workers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    2.2s
```

```
new size X_train (7995, 7)
new size X_train (3938, 7)
selected column names Index(['flushot6', 'hivrisk5', 'diffwalk', 'chcscncr',
'qstver', 'qstlang',
```

```

        'x.strwt'],
        dtype='object')

```

[Parallel(n_jobs=-1)]: Done 101 out of 101 | elapsed: 8.6s finished

[2022-12-12 21:06:38] Features: 7/7 -- score: 0.7333333333333333

```

[ ]: print("feature names", X_train_reduced4.columns)

```

```

feature names Index(['flushot6', 'hivrisk5', 'diffwalk', 'chcscncr', 'qstver',
                    'qstlang',
                    'x.strwt'],
                    dtype='object')

```

0.8.2 Model Creation : Initializing 5 models

Model 16 : Logistic Regression

```

[ ]: m16_lr, m16_y_pred_lr= lRMaker(X_train=X_train_reduced4, y_train=y_train,
    ↪X_test=X_test_reduced4)

```

```

[2 2 1 ... 2 2 1]

```

```

[ ]: m16_lr_result = printAccuracy(y_test, m16_y_pred_lr)

```

```

[2 2 1 ... 2 2 1]

```

```

      TN   FP   FN   TP  Accuracy  Sensitivity: TPR  Specificity: TNR
0  2444  858  186  450   0.734891           0.344037           0.929278
Log loss = 11.97184
AUC = 0.63666
Average Precision = 0.28472

```

Using 0.5 as threshold:

```


Accuracy = 0.73489
Precision = 0.70755
Recall = 0.34404
F1 score = 0.46296

```

Classification Report

	precision	recall	f1-score	support
1	0.71	0.34	0.46	1308
2	0.74	0.93	0.82	2630
accuracy			0.73	3938
macro avg	0.72	0.64	0.64	3938
weighted avg	0.73	0.73	0.70	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 17 :Decision Tree

```
[ ]: m17_dt, m17_y_pred_dt = dTMaker(X_train=X_train_reduced4, y_train=y_train,
    ↪X_test=X_test_reduced4)
```

```
[ ]: m17_dt_result = printAccuracy(y_test, m17_y_pred_dt)
```

```
[2 2 2 ... 2 2 1]
      TN  FP  FN  TP  Accuracy  Sensitivity: TPR  Specificity: TNR
0  2462  879  168  429  0.734129      0.327982      0.936122
Log loss = 11.97184
AUC = 0.63205
Average Precision = 0.28574
```

Using 0.5 as threshold:


```
Accuracy = 0.73413
Precision = 0.71859
Recall = 0.32798
F1 score = 0.45039
```

Classification Report

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

1	0.72	0.33	0.45	1308
2	0.74	0.94	0.82	2630
accuracy			0.73	3938
macro avg	0.73	0.63	0.64	3938
weighted avg	0.73	0.73	0.70	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 18 : KNN

```
[ ]: m18_knn, m18_y_pred_knn = kNNMaker(X_train=X_train_reduced4, y_train=y_train,
    ↪X_test=X_test_reduced4)
```

```
[ ]: m18_knn_result = printAccuracy(y_test, m18_y_pred_knn)
```

```
[2 2 2 ... 2 2 1]
      TN  FP  FN  TP  Accuracy  Sensitivity: TPR  Specificity: TNR
0  2298  772  332  536  0.719655      0.409786      0.873764
Log loss = 11.97184
AUC = 0.64178
Average Precision = 0.28453
```

Using 0.5 as threshold:

Accuracy = 0.71965

Precision = 0.61751


Recall = 0.40979

F1 score = 0.49265

Classification Report

	precision	recall	f1-score	support
1	0.62	0.41	0.49	1308
2	0.75	0.87	0.81	2630
accuracy			0.72	3938
macro avg	0.68	0.64	0.65	3938
weighted avg	0.71	0.72	0.70	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 19 : SVM

```
[ ]: m19_svm, m19_y_pred_svm = sVMMaker(X_train=X_train_reduced4, y_train=y_train,
↪X_test=X_test_reduced4)
```

```
[ ]: m19_svm_result = printAccuracy(y_test, m19_y_pred_svm)
```

```
[2 2 1 ... 2 2 1]
      TN   FP   FN   TP Accuracy Sensitivity: TPR Specificity: TNR
0  2425  836  205  472  0.735653          0.360856          0.922053
Log loss = 11.97184
AUC = 0.64145
Average Precision = 0.28371
```


Using 0.5 as threshold:

```
Accuracy = 0.73565
Precision = 0.69719
Recall = 0.36086
F1 score = 0.47557
```

Classification Report

	precision	recall	f1-score	support
1	0.70	0.36	0.48	1308
2	0.74	0.92	0.82	2630
accuracy			0.74	3938
macro avg	0.72	0.64	0.65	3938
weighted avg	0.73	0.74	0.71	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 20: MLP

```
[ ]: m20_mlp , m20_y_pred_mlp = mLPMaker(X_train=X_train_reduced4, y_train=y_train,
    ↪X_test=X_test_reduced4)
```

```
/home/harnold/github/2018_BRFSS_survey_data_anlaysia/env/lib/python3.8/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:679:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1) reached and the
optimization hasn't converged yet.
    warnings.warn(
```

```
[ ]: m20_mlp_result = printAccuracy(y_test, m20_y_pred_mlp)
```

```
[2 2 1 ... 2 2 1]
      TN    FP    FN    TP  Accuracy  Sensitivity: TPR  Specificity: TNR
0  2514  1044   116   264   0.705434          0.201835          0.955894
Log loss = 11.97184
AUC = 0.57886
Average Precision = 0.30124


Using 0.5 as threshold:
Accuracy = 0.70543
Precision = 0.69474
```

Recall = 0.20183
F1 score = 0.31280

Classification Report

	precision	recall	f1-score	support
1	0.69	0.20	0.31	1308
2	0.71	0.96	0.81	2630
accuracy			0.71	3938
macro avg	0.70	0.58	0.56	3938
weighted avg	0.70	0.71	0.65	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



0.9 Iteration 5

0.9.1 Attribute selection Method: Feature drop and manual elimination based on given info

```
[ ]: def dropFeatureAccr(x_train,x_test, y_train, y_test):
    accuracy_featuresDrop = []
    column_name= list(x_train.columns.values.tolist())
    for column in range(len(column_name)):
        x_train_modified = x_train.drop(columns=[column_name[column]])
        x_test_modified = x_test.drop(columns=[column_name[column]])
        log_reg_classifier = LogisticRegression(solver='lbfgs',
        ↪max_iter=2000,random_state=16)
        log_reg_classifier.fit(x_train_modified,y_train.values.ravel())
        y_pred_mod = log_reg_classifier.predict(x_test_modified)
        accuracy_featuresDrop.append(accuracy_score(y_test,y_pred_mod))
    dictionary_feature_drop = dict(zip(column_name, accuracy_featuresDrop))
    return dictionary_feature_drop

[ ]: logisticRegDrop =dropFeatureAccr(x_train= X_train_reduced1,y_train= y_train,
    ↪x_test= X_test_reduced1, y_test = y_test) #X_train_reduced1 , X_test_reduced1
print(logisticRegDrop)
print("max effect ")
print("feature", max(logisticRegDrop, key=logisticRegDrop.get))
print("accuracy score ", logisticRegDrop.get(max(logisticRegDrop,
    ↪key=logisticRegDrop.get)))
print("when we drop it the accurasy score went down effect")
print("feature",min(logisticRegDrop, key=logisticRegDrop.get))
print("accuracy score ", logisticRegDrop.get(min(logisticRegDrop,
    ↪key=logisticRegDrop.get)))

lists = logisticRegDrop.items() # sorted by key, return a list of tuples

x, y = zip(*lists) # unpack a list of pairs into two tuples


plt.plot(x, y)
plt.show()
```

```
{'x.age80': 0.6627729812087354, 'x.ageg5yr': 0.7148298628745556, 'x.age.g':
0.6945149822244794, 'employ1': 0.6937531742001015, 'diffwalk':
0.6945149822244794, 'x.hcvu651': 0.6950228542407313, 'genhlth':
0.6945149822244794, 'x.age65yr': 0.7158456069070595, 'marital':
0.6947689182326053, 'x.phys14d': 0.6947689182326053, 'physhlth':
0.6678517013712545, 'rmvteth4': 0.6947689182326053, 'pneuvac4':
0.7148298628745556, 'x.rfhlth': 0.6945149822244794, 'x.ststr':
0.7473336719146775, 'x.exteth3': 0.6945149822244794, 'persdoc2':
0.7150837988826816, 'qstver': 0.712798374809548, 'x.strwt': 0.6853732859319451,
```

```

'chccopd1': 0.6945149822244794}
max effect
feature x.ststr
accuracy score 0.7473336719146775
when we drop it the accuracy score went down effect
feature x.age80
accuracy score 0.6627729812087354

```



```

[ ]: features_dropped_box= ['x.ageg5yr', 'x.age65yr', 'x.ststr', 'x.strwt', 'x.age.g']
X_train_reduced5= X_train_reduced1.loc[ :, ~X_train_reduced1.columns.
    ↪isin(features_dropped_box)]

X_test_reduced5 = X_test_reduced1.loc[ :, ~X_test_reduced1.columns.
    ↪isin(features_dropped_box)]
X_train_reduced5.columns

[ ]: Index(['x.age80', 'employ1', 'diffwalk', 'x.hcvu651', 'genhlth', 'marital',
        'x.phys14d', 'physhlth', 'rmvteth4', 'pneuvac4', 'x.rfhlth',
        'x.exteth3', 'persdoc2', 'qstver', 'chccopd1'],
        dtype='object')

```

0.9.2 Model Creation : Initializing 5 models

Model 21 : LR

```
[ ]: m21_lr, m21_y_pred_lr= lRMaker(X_train = X_train_reduced5, y_train=y_train,  
↪X_test =X_test_reduced5)
```

```
[2 2 1 ... 2 2 1]
```

```
[ ]: m21_lr_result = printAccuracy(y_test, m21_y_pred_lr)
```

```
[2 2 1 ... 2 2 1]
```

	TN	FP	FN	TP	Accuracy	Sensitivity: TPR	Specificity: TNR
0	2276	647	354	661	0.74581	0.505352	0.865399

Log loss = 11.97184

AUC = 0.68538

Average Precision = 0.27734

Using 0.5 as threshold:

Accuracy = 0.74581

Precision = 0.65123


Recall = 0.50535

F1 score = 0.56909

Classification Report

	precision	recall	f1-score	support
1	0.65	0.51	0.57	1308
2	0.78	0.87	0.82	2630
accuracy			0.75	3938
macro avg	0.71	0.69	0.69	3938
weighted avg	0.74	0.75	0.74	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 22 : Decision Tree

```
[ ]: m22_dt, m22_y_pred_dt = dTMaker(X_train=X_train_reduced5, y_train=y_train,
    ↪X_test=X_test_reduced5)
```

```
[ ]: m22_lr_result = printAccuracy(y_test, m22_y_pred_dt)
```

```
[2 2 1 ... 2 2 1]
```

	TN	FP	FN	TP	Accuracy	Sensitivity: TPR	Specificity: TNR
0	2275	655	355	653	0.743525	0.499235	0.865019

Log loss = 11.97184

AUC = 0.68213

Average Precision = 0.27777

Using 0.5 as threshold:

Accuracy = 0.74352

Precision = 0.64782

Recall = 0.49924


F1 score = 0.56390

Classification Report

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

1	0.65	0.50	0.56	1308
2	0.78	0.87	0.82	2630
accuracy			0.74	3938
macro avg	0.71	0.68	0.69	3938
weighted avg	0.73	0.74	0.73	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 23 : KNN

```
[ ]: m23_knn, m23_y_pred_knn = kNNMaker(X_train=X_train_reduced5, y_train=y_train,
    ↪X_test=X_test_reduced5)
```

```
[ ]: m23_knn_result = printAccuracy(y_test, m23_y_pred_knn)
```

```
[2 2 1 ... 2 1 1]
      TN  FP  FN  TP Accuracy Sensitivity: TPR Specificity: TNR
0 2189 605 441 703 0.734383      0.537462      0.832319
Log loss = 11.97184
AUC = 0.68489
Average Precision = 0.27867
```

Using 0.5 as threshold:

Accuracy = 0.73438

Precision = 0.61451


Recall = 0.53746

F1 score = 0.57341

Classification Report

	precision	recall	f1-score	support
1	0.61	0.54	0.57	1308
2	0.78	0.83	0.81	2630
accuracy			0.73	3938
macro avg	0.70	0.68	0.69	3938
weighted avg	0.73	0.73	0.73	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 24 : SVM

```
[ ]: m24_svm, m24_y_pred_svm = sVMMaker(X_train=X_train_reduced5, y_train=y_train,
↪X_test=X_test_reduced5)
```

```
[ ]: m24_svm_result = printAccuracy(y_test, m24_y_pred_svm)
```

```
[2 2 1 ... 2 2 1]
```

	TN	FP	FN	TP	Accuracy	Sensitivity: TPR	Specificity: TNR
0	2345	685	285	623	0.753682	0.4763	0.891635

Log loss = 11.97184

AUC = 0.68397

Average Precision = 0.27660

Using 0.5 as threshold:

Accuracy = 0.75368

Precision = 0.68612


Recall = 0.47630

F1 score = 0.56227

Classification Report

	precision	recall	f1-score	support
1	0.69	0.48	0.56	1308
2	0.77	0.89	0.83	2630
accuracy			0.75	3938
macro avg	0.73	0.68	0.70	3938
weighted avg	0.74	0.75	0.74	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



Model 25 : MLP

```
[ ]: m25_mlp , m25_y_pred_mlp = mLPMaker(X_train=X_train_reduced5, y_train=y_train,
    ↪X_test=X_test_reduced5)
```

```
/home/harnold/github/2018_BRFSS_survey_data_anlaysia/env/lib/python3.8/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:679:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1) reached and the
optimization hasn't converged yet.
    warnings.warn(
```

```
[ ]: m25_mlp_result = printAccuracy(y_test, m25_y_pred_mlp)
```

```
[2 2 1 ... 2 2 1]
      TN   FP   FN   TP Accuracy Sensitivity: TPR Specificity: TNR
0  2317  670  313  638  0.750381      0.487768      0.880989
Log loss = 11.97184
AUC = 0.68438
Average Precision = 0.27691


Using 0.5 as threshold:
Accuracy = 0.75038
Precision = 0.67087
```


Recall = 0.48777
F1 score = 0.56485

Classification Report

	precision	recall	f1-score	support
1	0.67	0.49	0.56	1308
2	0.78	0.88	0.82	2630
accuracy			0.75	3938
macro avg	0.72	0.68	0.69	3938
weighted avg	0.74	0.75	0.74	3938

Note that in binary classification, recall of the positive class is also known as "sensitivity"; recall of the negative class is 'specificity'.



0.10 Saving Best Feature Data

```
[ ]: X_train_csv2= X_train_reduced5.copy()
X_train_csv2["havarth3"]= y_train.copy()
X_train_csv2.head()
X_train_csv2.to_csv("HicranArnold_best_train.csv", index=False)
```

```
X_test_csv2= X_test_reduced5.copy()
X_test_csv2["havarth3"]= y_test.copy()
X_test_csv2.to_csv("HicranArnold_best_test.csv", index=False)
```

0.11 Model Comparison

```
[ ]: y_pred_box = [m1_y_pred_lr, m2_y_pred_dt, m3_y_pred_knn, m4_y_pred_svm,
    ↪ m5_y_pred_mlp, m6_y_pred_lr, m7_y_pred_dt, m8_y_pred_knn, m9_y_pred_svm,
    ↪ m10_y_pred_mlp, m11_y_pred_lr, m12_y_pred_dt, m13_y_pred_knn,
    ↪ m14_y_pred_svm, m15_y_pred_mlp, m16_y_pred_lr, m17_y_pred_dt,
    ↪ m18_y_pred_knn, m19_y_pred_svm, m20_y_pred_mlp, m21_y_pred_lr,
    ↪ m22_y_pred_dt, m23_y_pred_knn, m24_y_pred_svm, m25_y_pred_mlp]
labels = ["m1", "m2", "m3", "m4", "m5", "m6", "m7", "m8", "m9", "m10", "m11",
    ↪ "m12", "m13", "m14", "m15", "m16", "m17", "m18", "m19", "m20", "m21", "m22",
    ↪ "m23", "m24", "m25"]

def accuracyRatesCollector(y_test, y_pred_box, labels ):
    df_box = []
    for item in range(len(y_pred_box)):
        y_pred= y_pred_box[item ]
        accuracy_rate = round(accuracy_score(y_test, y_pred),2)
        error_rate = 1- accuracy_rate
        sensitivity_recall= round(recall_score(y_test, y_pred),2)
        precision_rate = round(precision_score(y_test, y_pred),2)
        model_name = labels[item]
        my_matrix_values_1 = {"modelName":model_name, "accuracy":accuracy_rate,
    ↪ "error_rate":error_rate, "sensitivity_recall":
    ↪ sensitivity_recall, "precision_rate":precision_rate}
        df= pd.DataFrame(my_matrix_values_1 , index=[0])
        df_box.append(df)
    combined_df = pd.concat(df_box)
    return combined_df

[ ]: accuracy_mat = accuracyRatesCollector(y_test=y_test, y_pred_box=y_pred_box,
    ↪ labels= labels)
accuracy_mat
```

```
/home/harnold/github/2018_BRFSS_survey_data_anlaysia/env/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/home/harnold/github/2018_BRFSS_survey_data_anlaysia/env/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:
```

Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```


```
[ ]:  modelName  accuracy  error_rate  sensitivity_recall  precision_rate
0      m1      0.69      0.31      0.40      0.56
0      m2      0.74      0.26      0.50      0.65
0      m3      0.73      0.27      0.52      0.61
0      m4      0.75      0.25      0.47      0.69
0      m5      0.74      0.26      0.48      0.66
0      m6      0.67      0.33      0.00      0.00
0      m7      0.73      0.27      0.33      0.68
0      m8      0.72      0.28      0.48      0.60
0      m9      0.75      0.25      0.47      0.67
0     m10      0.74      0.26      0.50      0.65
0     m11      0.67      0.33      0.00      0.00
0     m12      0.74      0.26      0.50      0.65
0     m13      0.73      0.27      0.49      0.62
0     m14      0.75      0.25      0.47      0.68
0     m15      0.74      0.26      0.49      0.64
0     m16      0.73      0.27      0.34      0.71
0     m17      0.73      0.27      0.33      0.72
0     m18      0.72      0.28      0.41      0.62
0     m19      0.74      0.26      0.36      0.70
0     m20      0.71      0.29      0.20      0.69
0     m21      0.75      0.25      0.51      0.65
0     m22      0.74      0.26      0.50      0.65
0     m23      0.73      0.27      0.54      0.61
0     m24      0.75      0.25      0.48      0.69
0     m25      0.75      0.25      0.49      0.67
```

Plot Error Rate

```
[ ]: #accuracy_mat.T
```

```
[ ]: accuracy_mat.plot.bar(x='modelName', y='sensitivity_recall')
```

```
[ ]: <AxesSubplot: xlabel='modelName'>
```



```
[ ]: import matplotlib.pyplot as plt
import pandas as pd

# gca stands for 'get current axis'
ax = plt.gca()

accuracy_mat.plot(kind='line',x='modelName',y='sensitivity_recall',ax=ax)
accuracy_mat.plot(kind='line',x='modelName',y='precision_rate', color='red',ax=ax)

plt.show()
```

